

ОБЗОР SPECULUM - ЭТАЛОННАЯ ФОРМА ADMINISTRATUM-AGENT

Статус: ADVISORY / SPECULUM REVIEW ONLY

Коммит проверки: `d7bef1a60b0acbaec9fcde238711764f24fe00a6`

Репозиторий: <https://github.com/SoulsLike2313/Imperium->

Создано: 2026-05-18T01:40:10.901437Z

0. Объем и проверка

Проверяется Administratum-Agent V1 base в `IMPERIUM\_NEW\_GENERATION/ORGAN\_AGENTS/ADMINISTRATUM\_AGENT/`, а также target-form и Architectural Pattern Harvest. Локальный `git clone` из sandbox был невозможен из-за DNS-ошибки доступа к GitHub, поэтому runtime-тесты репозитория здесь не заявляются. HTTP GitHub-страницы открывались и использованы как источник фактов. Этот пакет - red-team и проектирование reference form, не acceptance receipt.

Что было прочитано

- `GitHub commit page` - Commit d7bef1a opens; message is TASK-20260518: implement Administratum Agent V1 base; page reports 71 files changed, +2510/-33.
- `IMPERIUM\_NEW\_GENERATION/README.md` - New Generation is a CLI-first sandbox; no canon merge, no UI/API integration; first 3 agents have common spine, 7 remain skeletons.
- `IMPERIUM\_NEW\_GENERATION/ORGAN\_AGENTS/ADMINISTRATUM\_AGENT/README.md` - Administratum-Agent mission covers inventory, classification, provenance, dirty runtime detection, useful candidates, merge intelligence, routing. Scope is sandbox-only, no canon fusion, no direct deletion/merge/promotion. Runtime outputs are under RUNS/ADMINISTRATUM_AGENT/<run_id>.
- `IMPERIUM\_NEW\_GENERATION/ORGAN\_AGENTS/ADMINISTRATUM\_AGENT/agent\_manifest.json` - Agent is IMPLEMENTED_V1_BASE, stdlib-only, RUNS_LAYER_ONLY, may write own state/memory/reports/receipts/outbox/runs, must not write canon files/other memory/core doctrine, and must not delete, merge, promote, bypass Custodes, ignore dirty state, or claim completion without receipts.
- `IMPERIUM\_NEW\_GENERATION/ORGAN\_AGENTS/ADMINISTRATUM\_AGENT/DESIGN\_TARGETS/ADMINISTRATUM\_AGENT\_V1\_TARGET\_FORM\_20260518/decision\_manifest.json` - Design target is TARGET_SPECIFICATION, RU/EN PDFs exist, OSS introduced none, future rich/Textual only considered and not installed/not canon; JSON receipts must remain unaffected.
- `IMPERIUM\_NEW\_GENERATION/RESEARCH/ARCHITECTURAL\_PATTERN\_HARVEST/TASK-IMPERIUM-OPEN-SOURCE-AGENT-PATTERN-HARVEST-V0\_1/README.md` - Pattern harvest is advisory research only; no external OSS framework is introduced, installed, or canonized. Key outputs include pattern catalog and native contract recommendations.
- `IMPERIUM\_NEW\_GENERATION/ORGAN\_AGENTS/ADMINISTRATUM\_AGENT/TOOLS/administratum\_v1\_core.py` - Core defines AGENT_ROOT, NEW_GENERATION_ROOT, REPO_ROOT, RUNS_ROOT, SKILL_IDS, UTF-8 JSON writes, classification with zone/artifact/provenance/trust/risk, route_for_object, inventory, dirty-runtime, routing and merge summary.
- `IMPERIUM\_NEW\_GENERATION/ORGAN\_AGENTS/ADMINISTRATUM\_AGENT/TOOLS/administratum\_agent\_runner.py` - Runner implements status, inventory, classify-path, provenance-index, detect-dirty-runtime, useful-candidates, route-to-organs, merge-summary, check-all; global flags are currently parser-level only, so flag placement after subcommand is not guaranteed.
- `IMPERIUM\_NEW\_GENERATION/ORGAN\_AGENTS/ADMINISTRATUM\_AGENT/TESTS/run\_check\_all.py` and runner check-all` - A check-all suite exists and validates manifest, policies, rules JSON, skill manifests, sample path classification, runtime classification, smoke inventory/provenance/routing/merge, CLI

no-color/plain-json, receipts, and no new dirty files outside allowed RUNS.

1. Жесткий вердикт

Administratum-Agent V1 - рабочая первая база. Это не идеальная форма и не эталон для копирования на остальные 9 агентов.

Уже есть правильный каркас: manifest, policies, brain_node, memory, skills, CLI runner, core logic, tests, receipts, reports, RUNS isolation. Но promotion в reference-grade заблокирован: нет исполняемой metrics layer, маршрутная логика слишком слабая, обучение не имеет admission-протокола, receipts неполные, fake-green semantics недостаточно жесткие, CLI UX не доведен до стандарта.

Следующий правильный шаг: сначала усилить форму, потом размножить агентов.

2. Что реально есть сейчас

Зона	Реальность V1	Red-team статус
Manifest	`agent_manifest.json` задает права, запреты, stdlib-only, skills.	Основа есть; нужна schema/gate проверка.
Policies	acceptance/rejection/learning/mutation/provenance/runtime policies есть.	Слишком коротко; еще не исполняемые gates.
Brain node	rules/cases/scoring/vocabulary/doctrine folders есть.	Нет unit accounting и memory hygiene.
Skills	7 skills и skill manifests.	Нужны schemas и fixture coverage.
CLI	status, inventory, classify, provenance, dirty runtime, candidates, routing, merge, check-all.	Полезно, но не reference-grade.
RUNS	outputs изолируются в RUNS/ADMINISTRATUM_AGENT.	Хорошо; gate должен блокировать загрязнение.
Tests	check-all существует.	Это smoke, не acceptance reference suite.
OSS	OSS не установлен и не принят; rich/Textual только future advisory.	Правильно, не ослаблять.

3. Главные блокеры и дыры

F-001 - Metrics - BLOCKER_FOR_REFERENCE_GRADE

Проблема: No first-class cost/time/resource layer exists. Reports do not yet standardize wall time, CPU time, peak memory, scanned/classified counts, output byte budget, warnings, rejects, route decisions, rerun cost, and human review estimate across every command.

Файлы: TOOLS/administratum_agent_runner.py, TOOLS/administratum_v1_core.py, state/current_status.json

Что сделать: Add AGENT_METRICS_V0_1 schema and write command_metrics.json, metrics section in every receipt, current_status.metrics, and run_index.jsonl per run.

Gate: metrics_gate must fail if any command lacks metrics fields.

F-002 - KPD usefulness - BLOCKER_FOR_REFERENCE_GRADE

Проблема: KPD review exists as a report artifact but no executable scoring model or testable KPD output is part of the agent runtime.

Файлы: REPORTS/TASK-20260518-BUILD-ADMINISTRATUM-AGENT-V1-FROM-TARGET-FORM/ADMINISTRATUM_AGENT_V1_KPD_REVIEW_V0_1.md, TOOLS/administratum_agent_runner.py

Что сделать: Define kpd_score.json and KPD formula with correctness, actionability, warning value, receipt completeness, false positives, false negatives, runtime cost, dirty-tree behavior, and repeated-work avoidance.

Gate: kpd_gate must require sample fixtures with expected score bands.

F-003 - Routing - BLOCKER_FOR_REFERENCE_GRADE

Проблема: route_for_object is too shallow. Default route is MECHANICUS_AGENT when no rule matches; runtime evidence and JSON reports can be routed without Administratum retention, Custodes risk, Inquisition suspicion, or source-context analysis.

Файлы: TOOLS/administratum_v1_core.py::route_for_object, brain_node/rules/routing_rules.json

Что сделать: Add routing_reason, route_confidence, route_policy_version, primary_owner, co_reviewers, escalation_reason, and fixtures for runtime JSON, receipts, policies, reports, archives, dirty outputs, and private-context references.

Gate: route_quality_gate must compare expected route sets for at least 50 fixtures.

F-004 - Learning safety - BLOCKER_FOR_REFERENCE_GRADE

Проблема: Learning memory directories exist, but there is no admission protocol that proves pending learning is reviewed, shadow-tested, admitted, quarantined, pruned, versioned, and reversible.

Файлы: memory/correction_memory/owner_corrections.jsonl, memory/learning_queue/pending_learning.jsonl, POLICIES/LEARNING_POLICY.md

Что сделать: Introduce LEARNING_ADMISSION_PROTOCOL_V0_1 with PENDING -> CANDIDATE -> SHADOW_VALIDATED -> OWNER_APPROVED -> ADMITTED -> DEPRECATED/QUARANTINED states and admission receipts.

Gate: learning_safety_gate must block any write to permanent memory without admission receipt.

F-005 - CLI UX - WARNING_HIGH

Проблема: CLI renderer is useful but not reference-grade. It lacks post-subcommand flag tolerance, compact mode, stable section order contract, severity legend, terminal width handling, and copy-friendly error pages.

Файлы: TOOLS/administratum_agent_runner.py::build_parser, DESIGN_TARGETS/.../decision_manifest.json

Что сделать: Adopt CLI_VISUAL_UX_STANDARD_V0_1. Flags must work before and after subcommands. Each command must emit header, verdict, warnings, outputs, receipts, metrics, next action, and JSON-equivalent plain mode.

Gate: cli_ux_gate must run command matrix in color/no-color/plain-json/verbose/compact modes on PowerShell-safe examples.

F-006 - Fake green - WARNING_HIGH

Проблема: status command can PASS because a snapshot was created, not because the agent is reference-grade. detect-dirty-runtime can return PASS_WITH_WARNINGS; if check-all treats warnings as

nonblocking, dirty signals may normalize as acceptable.

Файлы: TOOLS/administratum_agent_runner.py::command_status,
TOOLS/administratum_v1_core.py::detect_dirty_runtime_outputs

Что сделать: Split PASS, WARN, BLOCKED, FAIL semantics. Reference-grade gate must treat specific warnings as blockers depending on command mode and declared acceptance target.

Gate: no_fake_green_gate blocks promotion if any warning has blocker_class=true or unresolved_warning_count>0 for reference-grade claim.

F-007 - Receipt quality - WARNING_HIGH

Проблема: Receipts are generated but do not yet carry enough repeatability/provenance detail: command argv, cwd, environment hints, input SHA256, output SHA256, git HEAD, dirty state before/after, metrics, policy versions, and schema version are not consistently mandated.

Файлы: TOOLS/administratum_agent_runner.py receipts,
receipts/NEW_GENERATION_ADMINISTRATUM_AGENT_V1_IMPLEMENTATION_RECEIPT_V0_1.json

Что сделать: Define RECEIPT_SCHEMA_V0_2 and make every command use it. Keep old receipts as historical, not reference-grade.

Gate: receipt_completeness_gate validates every receipt with schema and hashes.

F-008 - Test depth - WARNING_MEDIUM

Проблема: check-all is a good V1 smoke gate but too narrow for reference-grade. It lacks golden fixtures, negative fixtures, corruption fixtures, route expectation fixtures, metric expectation fixtures, and repeated-run stability checks.

Файлы: TESTS/run_check_all.py, TOOLS/administratum_agent_runner.py::command_check_all

Что сделать: Add TEST_FIXTURES with expected outputs, expected warnings, and golden JSON snapshots. Add rerun stability test and output budget test.

Gate: test_depth_gate requires fixture coverage before reference promotion.

F-009 - State and dashboard readiness - WARNING_MEDIUM

Проблема: current_status exists but there is no strict dashboard-ready state model with last_checked_at, freshness window, stale status, last_run_id, last_receipt_path, last_blockers, KPD, metric summaries, and actionability.

Файлы: state/current_status.json, README.md runtime discipline

Что сделать: Define AGENT_CURRENT_STATUS_SCHEMA_V0_2 and dashboard_card.json. Dashboard must read only machine state and receipts, never console text.

Gate: dashboard_readiness_gate fails if stale state can be shown as green.

F-010 - 10k unit growth - WARNING_MEDIUM

Проблема: Brain-node directories are present, but there is no unit accounting model. The system cannot tell whether growth is useful knowledge or dirt accumulation.

Файлы: brain_node/rules, brain_node/cases, brain_node/scoring, brain_node/vocabulary, memory/*

Что сделать: Add AGENT_UNIT_INDEX.json with unit_id, category, source, status, confidence, last_used, usefulness_score, superseded_by, owner_reviewed, and pruning policy.

Gate: memory_hygiene_gate requires unit index and quarantine/prune receipts.

4. Что значит идеальная локальная CLI Organ-Agent форма

Локальный Organ-Agent должен быть маленьким, но серьезным мозговым узлом. Не LLM. Не облако. Не автономный владелец архитектуры. Его сила - в правилах, данных, маршрутах, receipts, metrics, controlled learning, reproducible outputs и честном отказе.

Минимальная форма:

- manifest с authority/read/write/forbidden actions;
- role contract;
- operating rules;
- policies;
- schemas;
- brain_node с индексируемыми units;
- memory с admission pipeline;
- skills со schema/fixtures;
- CLI runner/core;
- tests/golden/negative fixtures;
- receipts;
- current_status/dashboard_card;
- ignored RUNS layer.

5. Модель 10k meaningful units

10k units - это не веса нейросети. Это 10k осмысленных локальных единиц: rules, cases, route patterns, scoring weights, aliases, policy snippets, doctrine references, skill contracts, output templates, CLI display patterns, failure patterns, correction memory, fixtures, metrics, admission rules, risk signatures, accepted/rejected examples, Owner decisions, provenance examples.

Ключевой запрет: рост без admission/pruning - это не интеллект, это грязь.

Полная машинная модель: `AGENT\_10K\_UNIT\_MODEL.json`.

6. Cost/time/resource metrics

Каждая команда должна писать:

- wall-clock time;
- CPU time best-effort;

- peak memory best-effort;
- files scanned/classified;
- bytes read/written;
- report size;
- receipt count;
- warnings/blockers;
- rejected inputs;
- route decisions;
- useful candidates;
- dirty paths;
- human review estimate;
- rerun cost;
- task cost class.

Писать в: machine JSON, receipt metrics, human summary, current_status, run index.

7. KPD / коэффициент полезности

KPD должен быть вычисляемым, а не красивым отчетом. Базовая формула:

```
`KPD = clamp(0,100, 25*Correctness + 20*Actionability + 15*EvidenceCompleteness + 10*WarningValue + 10*NoFakeGreen + 10*NoDirtyTree + 10*Readability - 15*FalsePositiveRate - 25*FalseNegativeSeverity - 10*CostPenalty - 10*ManualCorrectionPenalty)`
```

Это V0.1, но уже годится как gate. Narrative KPD без JSON score - не gate.

8. Thinking quality для non-LLM агента

Мышление здесь - качество правил, данных, маршрутизации и scoring. Измерять нужно:

- classification confidence;
- route appropriateness;
- risk detection;
- evidence linking;
- provenance completeness;
- rejection correctness;
- contradiction awareness;
- unknown handling;
- escalation correctness;
- explanation quality;
- correction absorption;
- rerun stability.

9. CLI UX standard

Terminal должен быть приятным, но не в ущерб правде.

Обязательные правила:

- `--plain-json` = только JSON, без баннеров и ANSI;`
- `--no-color` = чистый PowerShell-safe output;`
- green только при evidence-backed PASS;
- WARN не равен PASS;
- поддерживать flags before/after subcommand;
- показывать header, run_id, verdict, blockers, outputs, receipts, metrics, next action;
- compact/verbose modes;
- UTF-8, без mojibake.

Полный стандарт: `CLI_VISUAL_UX_STANDARD.md`.`

10. Что нельзя копировать в остальные 9 агентов

Нельзя копировать	Почему	Замена
Default route to MECHANICUS	Скрывает неопределенность владельца.	UNKNOWN/NEEDS_REVIEW + owner/co-reviewers.
Narrative KPD	Не проверяется машиной.	<code>kpd_score.json`.</code>
Smoke check-all как proof	Покрытие слабое.	Golden/negative fixtures + rerun stability.
PASS за status snapshot	Fake green.	Status должен говорить readiness target и blockers.
PASS_WITH_WARNINGS для promotion	Warning blindness.	Warning severity policy.
Direct memory без admission	Memory dirt.	Learning admission state machine.
Только global flags	Операторская боль.	Flag placement tolerance.
Receipts без hashes/metrics	Слабый audit.	Receipt V0.2.

11. Что усиливать первым

1. Fake-green semantics и warning severity.
2. Receipt V0.2.
3. Metrics layer.
4. Route quality fixtures.
5. Learning admission.
6. KPD/self-review.
7. CLI UX.
8. Dashboard readiness.

Не начинай с визуального шлифования. Иначе будет красивая консоль, которая прячет слабые receipts и слабую маршрутизацию.

12. Roadmap

Administratum-Agent Hardening Roadmap - V1 Base to Reference Grade

Status: advisory only. Do not implement without a registered task.

Stage	Goal	Files affected	Tests needed	Outputs	Risk	Acceptance
Immediate fixes	Remove V1 friction and fake-green vectors	runner, check-all, status	flag placement, warning severity	flag receipt, no-fake-green receipt	parser churn	flags stable; warnings not hidden
V1.1 schemas	Standardize receipts/status/routes/errors	SCHEMAS, core, skills	schema fixtures	schema guard receipt	over-schema	every output validates
V1.2 metrics	Cost/time/resource layer	core, runner, status, run index	metrics presence	command_metric.s.json	Windows metric gaps	nulls require reason
V1.3 learning	Controlled correction memory	memory, policies, schemas	admission state machine	learning receipt, unit index	absorbing dirt	no permanent memory without admission
V1.4 KPD	Executable usefulness score	schemas/kpd, tools, fixtures	score bands	kpd_score.json	vanity scoring	reproducible score from evidence
V1.5 CLI UX	Pleasant terminal without truth loss	runner, CLI standard	mode matrix	cli_ux_receipt	polish hiding truth	plain-json clean; human output stable
V1.6 dashboard readiness	Dashboard-readable truth	state, dashboard_card, schemas	stale state fixtures	dashboard_card.json	dashboard source of truth	stale cannot show green
V2 reference gate	Promote reusable form	all gates	full gate suite	promotion candidate receipt	copying weak form	Owner accepts after gates

First hardening order

1. Fake-green and CLI flag placement.
2. Receipt schema V0.2.
3. Metrics model V0.1.
4. Route quality fixtures and routing confidence.
5. Learning admission protocol.
6. KPD/self-review.
7. CLI visual comfort.
8. Dashboard readiness.

Do not copy Administratum-Agent V1 into other agents before at least stages 1-4 are done.

13. Acceptance gates

Полная матрица: `REFERENCE\_GRADE\_ACCEPTANCE\_GATES.json`.

Критичные gates:

- truth gate;
- clean tree gate;

- runtime isolation gate;
- CLI UX gate;
- metrics gate;
- KPD gate;
- route quality gate;
- learning safety gate;
- memory hygiene gate;
- receipt completeness gate;
- no OSS dependency gate;
- no canon fusion gate.

14. Следующие Servitor-задачи

Task	Цель	Файлы	Acceptance
TASK-ADMIN-AGENT-V1_1-CLI-FAKEGREEN	Статусы, warnings, flags.	runner, tests	flags работают до/после subcommand; status не притворяется reference.
TASK-ADMIN-AGENT-V1_1-RECEIPT-SCHEMA	Receipt V0.2.	SCHEMAS, core, runner	receipts валидируются и содержат git/command/hashes/metrics.
TASK-ADMIN-AGENT-V1_2-METRICS	Metrics layer.	core, runner, state	каждая команда пишет metrics and cost class.
TASK-ADMIN-AGENT-V1_2-ROUTE-FIXTURES	Усилить routes.	rules, fixtures	runtime/report/policy/script/dirty routes проверены.
TASK-ADMIN-AGENT-V1_3-LEARNING-ADMISSION	Controlled learning.	memory, policies, schemas	permanent memory только через admission receipt.

15. Финальный вердикт

- Current V1 base: ACCEPTED AS WORKING BASE ONLY.
- Reference-grade: BLOCKED.
- Other 9 agents: DO NOT CLONE THIS FORM YET.
- OSS: NONE introduced; keep advisory only.
- Главные gaps: metrics, route quality, learning admission, receipt completeness, fake-green semantics.