

# Administratum-Agent V1 Target Form

Target architecture form for the first full Organ-Agent of IMPERIUM New Generation

**Document language**

English version for agents and Servitor

**Target placement**

IMPERIUM\_NEW\_GENERATION/ORGAN\_AGENTS/  
ADMINISTRATUM\_AGENT/DESIGN\_TARGETS/  
ADMINISTRATUM\_AGENT\_V1\_TARGET\_FORM\_20260518

**Reference HEAD**

8b4dd8eb218465a6c4fb07fbb80dadde26b40a7f

**Status**

TARGET SPECIFICATION NO FUSION

This document fixes the architecture, authority, prohibitions, folder form, memory/state/receipt model, CLI commands, readable terminal UX contract, and acceptance criteria for Administratum-Agent V1.

Generated for Owner review and later repository commit. Runtime output is not canon. This PDF is a target form, not an implementation receipt.

## Section index for future prompt references

| ID        | Section                            | ID        | Section                                 |
|-----------|------------------------------------|-----------|-----------------------------------------|
| ADM-V1-00 | Short definition                   | ADM-V1-16 | Memory Model                            |
| ADM-V1-01 | Agent mission                      | ADM-V1-17 | State Model                             |
| ADM-V1-02 | Place inside the IMPERIUM organism | ADM-V1-18 | CLI Commands                            |
| ADM-V1-03 | Permissions and prohibitions       | ADM-V1-19 | CLI Visual Comfort Contract             |
| ADM-V1-04 | Base agent folder                  | ADM-V1-20 | Receipt Contract                        |
| ADM-V1-05 | Why this structure exists          | ADM-V1-21 | Report Types                            |
| ADM-V1-06 | Agent manifest                     | ADM-V1-22 | Good / evil / learning                  |
| ADM-V1-07 | Role Contract                      | ADM-V1-23 | Performance / Optimization Rules        |
| ADM-V1-08 | Acceptance Policy                  | ADM-V1-24 | Interaction with Sanctum / Second Brain |
| ADM-V1-09 | Rejection Policy                   | ADM-V1-25 | Stress Tests                            |
| ADM-V1-10 | Learning Policy                    | ADM-V1-26 | V1 Acceptance Criteria                  |
| ADM-V1-11 | Mutation Policy and RUNS layer     | ADM-V1-27 | Large integration preparation task      |
| ADM-V1-12 | Repo Zone Classification           | ADM-V1-28 | First Build Plan                        |
| ADM-V1-13 | Artifact Classification            | ADM-V1-29 | Key architectural decision              |
| ADM-V1-14 | Provenance Index                   | ADM-V1-30 | Final form                              |
| ADM-V1-15 | Administratum-Agent V1 skills      | ADM-V1-31 | OSS / External Tooling Tracker          |

## Short definition

**Administratum-Agent** is the memory, archive, map, inventory, provenance, classification, and routing organ.

| Question                                  | Administratum-Agent answer |
|-------------------------------------------|----------------------------|
| Where is everything located?              | repo map / zone map        |
| Where did this object come from?          | provenance index           |
| What is canon, sandbox, runtime, or dirt? | classification engine      |
| Which scripts/gates/reports are useful?   | useful candidate scanner   |
| What must not be absorbed?                | rejection / dirt register  |
| Which organ should inspect it next?       | routing engine             |
| What should Servitor do?                  | merge preparation summary  |

**Main rule:** it does not integrate by itself. It prepares an exact map and summary so that Servitor can implement under Owner review.

## Agent mission

Administratum-Agent maintains the operational memory, archive, repository map, provenance index, artifact classification, routing recommendations, and merge preparation summaries of IMPERIUM.

It inspects, indexes, classifies, summarizes, routes, and recommends.  
It does not mutate canon, delete files, silently promote artifacts, or perform uncontrolled merges.

Plainly: **Administratum-Agent is the archivist, cartographer, and dispatcher of IMPERIUM.** It knows what is located where, where it came from, what its status is, whether it may be absorbed, which organ should review it, and what Servitor should do next.

## Place inside the IMPERIUM organism

```

Owner / Servitor / Sanctum
↓
Task Envelope / Repo Snapshot / Artifact Packet
↓
ADMINISTRATUM_AGENT
↓
Repo Map
Provenance Index
Classification Report
Useful Candidates
Dirty / Forbidden Items
Routing Recommendations
Merge Preparation Summary
↓
Other Organ-Agents
↓
Servitor Implementation
↓
Receipts / Gates / Clean Tree

```

Administratum is the first understanding filter. It does not solve everything, but without it no organ knows where truth, dirt, old canon, runtime, and sandbox boundaries are.

## Permissions and prohibitions

| Action                        | Allowed? | Comment                                     |
|-------------------------------|----------|---------------------------------------------|
| Read repo snapshot            | Yes      | through index/scanner                       |
| Build file map                | Yes      | without canon mutation                      |
| Classify zones                | Yes      | canon/sandbox/runtime/private/etc           |
| Create reports                | Yes      | inside its own outbox/reports/runtime layer |
| Create receipts               | Yes      | mandatory evidence                          |
| Recommend actions to Servitor | Yes      | but not execute merge                       |
| Update own memory             | Yes      | only in own zone or runtime/state           |
| Write canon repo directly     | No       | only through Servitor task                  |
| Delete files                  | No       | even if marked as dirt                      |
| Merge old and new systems     | No       | summary only for Owner/Servitor             |
| Declare artifact canon        | No       | candidate/recommended only                  |
| Ignore dirty state            | No       | must report it                              |

## Base agent folder

---

```
IMPERIUM_NEW_GENERATION/  
  ORGAN_AGENTS/  
    ADMINISTRATUM_AGENT/  
      AGENT_MANIFEST.json  
      ROLE_CONTRACT.md  
      README.md  
  
    POLICIES/  
      ACCEPTANCE_POLICY.md  
      REJECTION_POLICY.md  
      LEARNING_POLICY.md  
      MUTATION_POLICY.md  
      PROVENANCE_POLICY.md  
      RUNTIME_OUTPUT_POLICY.md  
  
    BRAIN_NODE/  
      rules/  
        repo_zone_classification_rules.json  
        artifact_acceptance_rules.json  
        artifact_rejection_rules.json  
        routing_rules.json  
        merge_candidate_rules.json  
        dirt_detection_rules.json  
      cases/  
        known_good_cases.jsonl  
        known_bad_cases.jsonl  
        correction_cases.jsonl  
        owner_decision_cases.jsonl  
      vocabulary/  
        canonical_terms.json  
        aliases.json  
        forbidden_terms.json  
        path_markers.json  
      scoring/  
        relevance_weights.json  
        trust_weights.json  
        freshness_weights.json  
        merge_value_weights.json  
        risk_weights.json  
      doctrine/  
        administratum_doctrine.md  
        archive_doctrine.md  
        repo_truth_doctrine.md  
  
    SKILLS/  
      build_repo_inventory/  
      classify_repo_zone/  
      build_provenance_index/  
      find_useful_candidates/  
      detect_dirty_runtime_outputs/  
      build_merge_preparation_summary/  
      route_to_organs/  
  
    MEMORY/  
      permanent/  
      correction_memory/  
      learning_queue/  
  
    INBOX/  
      task_envelopes/  
      repo_snapshots/  
      artifact_packets/  
      correction_packets/  
      organ_reports/  
  
    OUTBOX/  
      answers/  
      inventory_summaries/
```

```
merge_preparation_summaries/  
routing_recommendations/  
rejection_reports/
```

```
STATE/  
RECEIPTS/  
REPORTS/  
TESTS/  
TOOLS/
```

#### ADM-V1-05

## Why this structure exists

| Folder                | Purpose                           |
|-----------------------|-----------------------------------|
| POLICIES              | human-readable agent laws         |
| BRAIN_NODE/rules      | machine-readable rules            |
| BRAIN_NODE/cases      | known good/bad/corrected examples |
| BRAIN_NODE/vocabulary | path/term/alias dictionary        |
| BRAIN_NODE/scoring    | weights for value and risk        |
| SKILLS                | executable capabilities           |
| MEMORY                | long-term agent memory            |
| INBOX                 | incoming tasks and packets        |
| OUTBOX                | answers leaving the agent         |
| STATE                 | current agent state               |
| RECEIPTS              | evidence of actions               |
| REPORTS               | analytical reports                |
| TESTS                 | behavior verification             |
| TOOLS                 | local runner/validator            |

Key principle: **everything the agent knows, can do, accepts, rejects, and performs must be visible in its folder structure.**

#### ADM-V1-06

## Agent manifest

`AGENT_MANIFEST.json` is the fast machine-readable passport of the agent.

```
{
  "agent_id": "ADMINISTRATUM_AGENT",
  "agent_name": "Administratum-Agent",
  "version": "V1_DRAFT",
  "organ": "ADMINISTRATUM",
  "status": "SPECIFICATION_DRAFT",
  "primary_role": "memory_archive_inventory_provenance_routing",
  "may_read": ["repo_snapshots", "task_envelopes", "artifact_packets", "organ_reports",
"continuity_packs"],
  "may_write": ["own_state", "own_memory", "own_reports", "own_receipts", "own_outbox"],
  "must_not_write": ["canon_repo_files", "other_agent_memory", "core_doctrine",
"task_implementation_files"],
  "must_not_do": ["delete_files", "merge_artifacts", "promote_to_canon", "bypass_custodes",
"ignore_dirty_state", "claim_completion_without_receipts"],
  "required_receipts": true,
  "owner_review_required_for": ["canon_promotion", "deletion_recommendations",
"merge_recommendations", "unknown_high_risk_artifacts"]
}
```

ADM-V1-07

## Role Contract

ROLE\_CONTRACT.md must be strict, not decorative. It fixes identity, mission, authority, limits, inputs, outputs, receipts, escalation, failure behavior, and learning.

Administratum may classify and recommend.  
Administratum must not mutate or merge.

ADM-V1-08

## Acceptance Policy

| Input type       | Acceptance condition                                       |
|------------------|------------------------------------------------------------|
| TASK_ENVELOPE    | has task_id, source, intent, allowed scope                 |
| REPO_SNAPSHOT    | has root, git_head, status, file index or tree             |
| ARTIFACT_PACKET  | has path/source/provenance or is explicitly marked unknown |
| ORGAN_REPORT     | has agent_id, report_type, timestamp/run_id                |
| OWNER_CORRECTION | has correction target and human note                       |
| CONTINUITY_PACK  | has bundle id or source path                               |

| Rejection reason                 | Action                           |
|----------------------------------|----------------------------------|
| missing source/provenance        | reject or mark UNKNOWN           |
| task asks to delete files        | escalate to Custodes/Owner       |
| task asks to merge               | create preparation summary only  |
| runtime artifact without receipt | mark RUNTIME_UNTRUSTED           |
| unknown private/local path       | mark PRIVATE_OR_EXTERNAL_UNKNOWN |
| path outside allowed scope       | reject + route to Custodes       |



## Rejection Policy

| Verdict                       | Meaning                                                     |
|-------------------------------|-------------------------------------------------------------|
| REJECT_MISSING_PROVENANCE     | origin is unclear                                           |
| REJECT_FORBIDDEN_SCOPE        | object is outside allowed scope                             |
| REJECT_MUTATION_REQUEST       | task asks for an action the agent is not allowed to perform |
| REJECT_RUNTIME_AS_CANON       | runtime is being treated as canon                           |
| REJECT_UNSAFE_PRIVATE_CONTEXT | private/local context has no admission rules                |
| REJECT_FAKE_GREEN_RISK        | result has no receipts/checks                               |
| REJECT_UNKNOWN_HIGH_RISK      | too risky without Owner review                              |

Every rejection must write a receipt.

## Learning Policy

| Source                   | May learn?   | How                             |
|--------------------------|--------------|---------------------------------|
| Owner correction         | Yes          | directly into correction memory |
| Verified Servitor result | Yes          | after receipt/check             |
| Passed stress test       | Yes          | as known good case              |
| Failed classification    | Yes          | as failed case                  |
| Advisory text            | Carefully    | advisory only, not canon        |
| Runtime output           | Not directly | only after classification       |
| Dirty state              | Not as truth | only as dirty case              |
| Internet/OSS note        | Not directly | through advisory/review         |

No learning without source, classification, and receipt.

## Mutation Policy and RUNS layer

Administratum V1 must not write into the main repository except inside its permitted own zone or runtime layer.

| Zone                       | May write? |
|----------------------------|------------|
| ADMINISTRATUM_AGENT/STATE  | Yes        |
| ADMINISTRATUM_AGENT/MEMORY | Yes        |

| Zone                         | May write?                     |
|------------------------------|--------------------------------|
| ADMINISTRATUM_AGENT/REPORTS  | Yes                            |
| ADMINISTRATUM_AGENT/RECEIPTS | Yes                            |
| ADMINISTRATUM_AGENT/OUTBOX   | Yes                            |
| ORGANS/... canon             | No                             |
| IMPERIUM_TEST_VERSION/...    | No                             |
| other agent folder           | No                             |
| root files                   | No                             |
| tracked runtime outputs      | Undesired / fix via RUNS layer |

```

IMPERIUM_NEW_GENERATION/
  RUNS/
    .gitkeep
    ADMINISTRATUM_AGENT/
      RUN-20260518-0001/
        input_snapshot.json
        inventory_report.json
        classification_report.json
        useful_candidates.json
        receipt.json
        state_delta.json

```

#### ADM-V1-12

## Repo Zone Classification

| Zone class             | Meaning                              |
|------------------------|--------------------------------------|
| CANON_CORE             | main IMPERIUM canon                  |
| ORGAN_CANON            | organ canon zone                     |
| SANCTUM_UI             | Sanctum/dashboard/UI                 |
| SECOND_BRAIN           | Second Brain / operator brain        |
| TEST_VERSION           | experimental test version            |
| NEW_GENERATION_SANDBOX | new agent sandbox zone               |
| RUNTIME_OUTPUT         | execution output                     |
| RECEIPT                | evidence                             |
| REPORT                 | report                               |
| SCRIPT                 | executable script                    |
| TOOL                   | tool/adapter                         |
| POLICY                 | rule/policy                          |
| DOCTRINE               | law/doctrine                         |
| ADVISORY               | Kiro/Speculum/Deep Research advisory |

| Zone class     | Meaning                      |
|----------------|------------------------------|
| LOCAL_PRIVATE  | private/local context        |
| ARCHIVE_LEGACY | old/archive material         |
| DIRTY_UNKNOWN  | unclear pollution            |
| FORBIDDEN      | must not be absorbed/touched |

#### ADM-V1-13

## Artifact Classification

| Artifact type  | Example                |
|----------------|------------------------|
| SCRIPT         | .py, .ps1, .sh         |
| SCHEMA         | .schema.json           |
| POLICY_DOC     | policy markdown        |
| DOCTRINE_DOC   | law/doctrine markdown  |
| REPORT_JSON    | generated report       |
| RECEIPT_JSON   | run receipt            |
| DASHBOARD_DATA | dashboard JSON         |
| UI_FILE        | html/css/js            |
| ADVISORY_DOC   | Kiro/Speculum/research |
| TASK_DRAFT     | task/stage draft       |
| RUNTIME_STATE  | state after run        |
| UNKNOWN        | unclear                |

```
{
  "artifact_id": "sha256_or_path_based_id",
  "path": "IMPERIUM_NEW_GENERATION/TOOLS/agent_cli/agent_cli.py",
  "zone_class": "NEW_GENERATION_SANDBOX",
  "artifact_type": "SCRIPT",
  "provenance_status": "KNOWN_COMMITTED",
  "trust_level": "SANDBOX_TRUSTED_IF_CLEAN",
  "merge_value": "POTENTIALLY_USEFUL",
  "risk_level": "LOW",
  "recommended_next_route": ["MECHANICUS_AGENT"],
  "owner_review_required": false
}
```

#### ADM-V1-14

## Provenance Index

| Field | Meaning          |
|-------|------------------|
| path  | file/object path |

| Field               | Meaning                                  |
|---------------------|------------------------------------------|
| git_head_seen       | commit/head where object was seen        |
| first_seen          | first known moment                       |
| source_type         | commit, bundle, local, advisory, runtime |
| source_id           | commit sha / bundle id / run id          |
| created_by          | Owner, Servitor, Kiro, script, unknown   |
| verified_by         | checker/receipt/manual                   |
| trust_status        | trusted/advisory/unknown/dirty           |
| last_classification | latest classification                    |
| notes               | short explanation                        |

#### ADM-V1-15

## Administratum-Agent V1 skills

| Skill                           | Purpose                                     | Output                   |
|---------------------------------|---------------------------------------------|--------------------------|
| build_repo_inventory            | build repository map                        | inventory report         |
| classify_repo_zone              | classify path/object zone                   | zone classification card |
| build_provenance_index          | collect object origin data                  | provenance report        |
| find_useful_candidates          | find useful items for agent-era integration | useful candidates report |
| detect_dirty_runtime_outputs    | detect runtime/state pollution              | dirty runtime report     |
| build_merge_preparation_summary | create Owner/Servitor summary               | merge preparation packet |
| route_to_organs                 | decide which organ should review next       | routing recommendations  |

Useful candidates: gates/checkers, script registries, dashboards, policies, receipts, task/stage schemas, performance tools, Kiro/Speculum reports. Do not absorb directly: runtime dirty outputs, broken UI experiments, generated reports without source, private/local dumps, duplicate stale docs.

| Finding                                | Route                 |
|----------------------------------------|-----------------------|
| script/tool/checker                    | MECHANICUS_AGENT      |
| fake green / duplicate / contradiction | INQUISITION_AGENT     |
| dangerous path / permission / deletion | CUSTODES_AGENT        |
| policy/law/doctrine                    | DOCTRINARIUM_AGENT    |
| task/stage/roadmap                     | ASTRONOMICON_AGENT    |
| role/agent behavior                    | OFFICIO_AGENTIS_AGENT |

| Finding                      | Route                   |
|------------------------------|-------------------------|
| learning/curriculum/examples | SCHOLA_IMPERIALIS_AGENT |
| strategy/priority            | STRATEGIUM_AGENT        |
| supreme decision             | THRONE_AGENT            |

#### ADM-V1-16

## Memory Model

Administratum memory must not become a pile of prose. Use structured JSONL.

```
repo_map_memory.jsonl
{"event_type": "REPO_OBJECT_CLASSIFIED", "path": "...", "zone_class": "NEW_GENERATION_SANDBOX", "artifact_type": "S"}

decision_memory.jsonl
{"event_type": "OWNER_DECISION", "decision_id": "DECISION-NEWGEN-NO-FUSION-YET", "summary": "Do not fuse New Generation after first 3 agents. Build all 10 first.", "scope": "IMPERIUM_NEW_GENERATION", "status": "ACTIVE"}

owner_corrections.jsonl
{"correction_id": "CORR-0001", "target": "tracked_runtime_outputs", "owner_note": "Runtime outputs prove execution but must not remain dirty in tracked tree.", "new_rule": "Move runtime outputs to RUNS or restore after manual run.", "status": "ACTIVE"}
```

#### ADM-V1-17

## State Model

```
{
  "agent_id": "ADMINISTRATUM_AGENT",
  "status": "READY_FOR_CLASSIFICATION_TASKS",
  "last_run_id": "RUN-ADMINISTRATUM-0001",
  "known_limits": [
    "Does not mutate canon",
    "Does not delete files",
    "Does not perform merges",
    "Requires provenance for canon recommendations"
  ],
  "capabilities": {
    "repo_inventory": "AVAILABLE",
    "zone_classification": "AVAILABLE",
    "provenance_indexing": "DRAFT",
    "useful_candidate_scan": "DRAFT",
    "merge_preparation_summary": "DRAFT"
  }
}
```

#### ADM-V1-18

## CLI Commands

Administratum must expose understandable commands. A command must not only print something nice. It must write machine files: `report.json`, `summary.md`, `receipt.json`, and `state_delta.json`.

```
python TOOLS/administratum_agent_runner.py status

python TOOLS/administratum_agent_runner.py inventory --repo-root E:\IMPERIUM --git-head
8b4dd8eb218465a6c4fb07fbb80dadde26b40a7f --out RUNS/ADMINISTRATUM_AGENT/<run_id>/

python TOOLS/administratum_agent_runner.py classify-path --path IMPERIUM_NEW_GENERATION/TOOLS/
agent_cli/runner.py

python TOOLS/administratum_agent_runner.py detect-dirty-runtime --repo-root E:\IMPERIUM

python TOOLS/administratum_agent_runner.py useful-candidates --inventory RUNS/
ADMINISTRATUM_AGENT/<run_id>/inventory.json

python TOOLS/administratum_agent_runner.py merge-summary --inventory ... --provenance ...
--out RUNS/ADMINISTRATUM_AGENT/<run_id>/MERGE_PREPARATION_SUMMARY/
```

## ADM-V1-19

# CLI Visual Comfort Contract

**New Owner requirement:** CLI must not be a wall of white lines. It must be a readable local work surface. Beauty is allowed only as a clarity layer. It must never hide warnings, blockers, receipts, or fake-green risks.

| Rule                    | Requirement                                                                               |
|-------------------------|-------------------------------------------------------------------------------------------|
| Readable sections       | each command prints short colored headers, verdict, outputs, next action                  |
| Color discipline        | cyan headers, green PASS, yellow WARN, red BLOCK/FAIL, gray paths/metadata                |
| ASCII-safe fallback     | no mandatory glyphs that break in PowerShell/encoding; color is more important than emoji |
| No fake green           | green PASS only when it is a real PASS; WARN/BLOCK must remain visible                    |
| Machine output separate | pretty terminal output never replaces JSON/MD/receipt files                               |
| Flags                   | required modes: --no-color, --plain, --json, --quiet                                      |
| Compact tables          | paths, statuses, counts shown as tables/cards, not endless prose                          |
| Output budget           | long lists go to report files; CLI prints summary + paths                                 |
| Operator pleasure       | local runs should be pleasant to read without AAA UI or bloat                             |
| Log safety              | color/beauty must not break plain logs, CI, grep, receipts                                |

```

=== ADMINISTRATUM AGENT :: INVENTORY RUN ===
run_id RUN-ADMINISTRATUM-20260518-0001
repo E:\IMPERIUM
head 8b4dd8e

PASS repo snapshot accepted
WARN runtime outputs detected in tracked New Generation paths
BLOCK no canon mutation allowed from this command

Outputs
report: RUNS/ADMINISTRATUM_AGENT/RUN-.../inventory_report.json
summary: RUNS/ADMINISTRATUM_AGENT/RUN-.../summary.md
receipt: RUNS/ADMINISTRATUM_AGENT/RUN-.../receipt.json

Verdict PASS_WITH_WARNINGS
Next backup patch, restore tracked runtime, design RUNS layer

```

A future implementation may consider `rich` or a similar local renderer for Python CLI, but this PDF does not install anything and does not promote such dependency to canon.

ADM-V1-20

## Receipt Contract

```

{
  "receipt_type": "ADMINISTRATUM_SKILL_RUN_RECEIPT",
  "agent_id": "ADMINISTRATUM_AGENT",
  "run_id": "RUN-ADMINISTRATUM-20260518-0001",
  "skill_id": "build_repo_inventory",
  "input_refs": [
    "repo_root:E:/IMPERIUM",
    "git_head:8b4dd8eb218465a6c4fb07fbb80dadde26b40a7f"
  ],
  "outputs": ["inventory_report.json", "inventory_summary.md"],
  "verdict": "PASS",
  "warnings": [],
  "mutated_canon": false,
  "deleted_files": false,
  "timestamp": "2026-05-18T00:00:00Z"
}

```

| Field         | Required |
|---------------|----------|
| agent_id      | Yes      |
| run_id        | Yes      |
| skill_id      | Yes      |
| input_refs    | Yes      |
| outputs       | Yes      |
| verdict       | Yes      |
| mutated_canon | Yes      |
| deleted_files | Yes      |
| warnings      | Yes      |

## Report Types

```

REPO_INVENTORY_REPORT
  git_head
  total_files
  zone_counts
  top_level_dirs
  objects[]

ARTIFACT_CLASSIFICATION_REPORT
  classified_count
  unknown_count
  dirty_count
  forbidden_count
  high_risk_count

USEFUL_CANDIDATES_REPORT
  script_candidates
  gate_candidates
  policy_candidates
  dashboard_candidates
  memory_candidates
  requires_mechanicus_review
  requires_inquisition_review
  requires_owner_decision

DIRTY_RUNTIME_REPORT
  dirty_runtime_detected
  dirty_paths
  recommended_action
  architectural_fix

```

## Good / evil / learning

| Category | Meaning                                            |
|----------|----------------------------------------------------|
| Good     | clean map, clear provenance, honest classification |
| Good     | useful summary for Owner/Servitor                  |
| Good     | canon/sandbox/runtime/private separation           |
| Good     | receipts for every action                          |
| Good     | route to the right organ                           |
| Evil     | fake green                                         |
| Evil     | runtime as canon                                   |
| Evil     | unknown source promoted as truth                   |
| Evil     | merge without Owner/Servitor path                  |
| Evil     | deletion without Custodes/Owner                    |
| Evil     | latest without hash/provenance                     |
| Learning | Owner correction                                   |



| Category     | Meaning                         |
|--------------|---------------------------------|
| Learning     | verified Servitor result        |
| Learning     | failed case turned into a rule  |
| Learning     | stress test result with receipt |
| Not learning | unverified advisory             |
| Not learning | dirty runtime                   |
| Not learning | random generated file           |

ADM-V1-23

## Performance / Optimization Rules

| Risk                      | Solution                                    |
|---------------------------|---------------------------------------------|
| Full repo scan every time | incremental inventory                       |
| Large reports             | report output budget                        |
| Huge JSON                 | JSONL + chunked reports                     |
| UI lag                    | CLI writes state, Sanctum only reads        |
| repeated classification   | cache by path + hash                        |
| runtime dirty             | RUNS layer ignored by Git                   |
| memory bloat              | memory compaction reports                   |
| long command blocks       | chunked execution                           |
| Windows command issues    | small CLI calls, no giant PowerShell blocks |
| fake pass                 | receipts + explicit warnings                |

Administratum V1 should be boring and fast. CLI beauty is permitted as readability, not decorative lying.

## Interaction with Sanctum / Second Brain

```

Sanctum / Second Brain UI
reads:
  STATE/current_status.json
  REPORTS/latest_summary.json
  RECEIPTS/latest_receipt.json
  LEDGER/events.jsonl

triggers:
  CLI command / task envelope

does not:
  scan repo itself
  decide canon status itself
  mutate agent memory directly

```

| Panel                  | Data source                   |
|------------------------|-------------------------------|
| Repo zones             | inventory report              |
| Dirty runtime warning  | dirty runtime report          |
| Useful candidates      | useful candidates report      |
| Agent status           | current_status.json           |
| Latest receipts        | receipts index                |
| Owner decisions needed | pending_owner_decisions.jsonl |
| Route recommendations  | routing report                |

## Stress Tests

| Test                            | Expected verdict           |
|---------------------------------|----------------------------|
| clean repo snapshot             | PASS_INVENTORY_BUILT       |
| path in IMPERIUM_NEW_GENERATION | NEW_GENERATION_SANDBOX     |
| path in RUNS                    | RUNTIME_OUTPUT             |
| receipt file                    | RECEIPT                    |
| advisory file                   | ADVISORY                   |
| unknown untracked file          | DIRTY_UNKNOWN              |
| task asks to delete             | REJECT_MUTATION_REQUEST    |
| task asks to merge              | PREPARE_SUMMARY_ONLY       |
| runtime result as canon         | REJECT_RUNTIME_AS_CANON    |
| script candidate found          | route to MECHANICUS_AGENT  |
| duplicate/fake-green suspicion  | route to INQUISITION_AGENT |

## V1 Acceptance Criteria

| Criterion                                   | Required status                     |
|---------------------------------------------|-------------------------------------|
| Manifest exists                             | PASS                                |
| Role contract exists                        | PASS                                |
| Acceptance/rejection policies exist         | PASS                                |
| Zone classification rules exist             | PASS                                |
| Inventory skill works                       | PASS                                |
| Dirty runtime detector works                | PASS                                |
| Useful candidate report is created          | PASS                                |
| Merge preparation summary is created        | PASS                                |
| Receipts are written                        | PASS                                |
| No canon mutation                           | PASS                                |
| No delete                                   | PASS                                |
| Tests pass                                  | PASS                                |
| Output is controlled                        | PASS                                |
| Runtime does not dirty tracked architecture | PASS or known blocker with fix plan |
| CLI readable and pleasant                   | PASS_WITH_NO_TRUTH_LOSS             |

## Large integration preparation task

```
ADMINISTRATUM_MERGE_PREPARATION_SUMMARY
1. Current repo zone map
2. Canon memory candidates
3. Useful scripts and gates
4. Useful dashboards and state models
5. Useful policies/doctrine
6. Runtime/dirty items that must not be absorbed
7. Deprecated/legacy items
8. Unknown items needing review
9. Routes to other Organ-Agents
10. Recommended Servitor action list
11. Owner decision list
12. Receipts and evidence
```

The agent must not silently absorb all useful material. It must say: this is useful, this is why, this is where it lives, this is its provenance, this is the risk, this is who should review it, this is what Servitor should do.

## First Build Plan

| Stage | Name                      | Result                                       |
|-------|---------------------------|----------------------------------------------|
| 1     | Identity & Policy Base    | manifest, role contract, policies            |
| 2     | Zone Classification Spine | zone rules, classifier, fixtures             |
| 3     | Inventory Builder         | repo inventory skill + report                |
| 4     | Provenance Index Base     | provenance schema + index report             |
| 5     | Dirty Runtime Detector    | detector for runtime/state pollution         |
| 6     | Useful Candidate Scanner  | scripts/gates/policies/dashboard candidates  |
| 7     | Routing Engine            | route recommendations to organs              |
| 8     | Merge Preparation Summary | Owner/Servitor summary format                |
| 9     | Receipts & State          | run receipts, skill receipts, current_status |
| 10    | Stress Suite              | acceptance/rejection/performance tests       |

## Key architectural decision

```
ARCHITECTURE FILES
  committed, stable, reviewed

RUNTIME OUTPUTS
  generated by runs, ignored or stored under RUNS

MEMORY
  controlled, append-only or receipt-backed

REPORTS
  generated evidence, not canon by default

CANON PROMOTION
  only via Owner/Servitor/gates
```

Otherwise every agent run will dirty Git and require manual restoration of tracked runtime files again.

## Final form

Input:  
task envelope / repo snapshot / artifact packet / correction

Process:  
accept/reject  
classify  
index  
score  
route  
summarize  
receipt

Output:  
reports  
recommendations  
merge preparation summary  
receipts  
state update

Never:  
delete  
merge  
mutate canon  
promote unknown to truth  
hide dirty state

| Agent           | Main meaning                                        |
|-----------------|-----------------------------------------------------|
| Custodes        | permission / protection / forbidden access          |
| Inquisition     | contradiction / fake green / duplicate / corruption |
| Mechanicus      | scripts / tools / execution machinery               |
| Doctrinarium    | law / canon / doctrine                              |
| Astronomicon    | tasks / stages / roadmaps                           |
| Officio Agentis | roles / behavior / response contracts               |
| Strategium      | strategy / priorities / campaigns                   |
| Schola          | learning / examples / training                      |
| Throne          | final authority / imperial decision layer           |

## OSS / External Tooling Tracker

| Field                                  | Status                                               |
|----------------------------------------|------------------------------------------------------|
| OSS introduced now                     | NONE                                                 |
| OSS considered for future CLI renderer | Python rich / Textual-style renderer - advisory only |
| Install status                         | NOT INSTALLED BY THIS DOCUMENT                       |

| Field                                 | Status                                                                                               |
|---------------------------------------|------------------------------------------------------------------------------------------------------|
| Canon status                          | NOT CANON                                                                                            |
| Why considered                        | readable colored local CLI without fake green                                                        |
| Verification required before adoption | local install receipt, no mojibake, plain fallback, JSON receipts unaffected, acceptable performance |